

Corporate workshop • Half day • Module A10

AI in your close, checks built in

You leave with three pieces of your own close work drafted by AI, each one tied out against a number the model did not produce.

Varun Tyagi • Raygency

September 27, 2026 • raygency.com

How this runs

0:00

Your newest preparer has never seen your ledger

What the model is doing when it writes your commentary, and where it is safe to start

0:55

Break

Ten minutes. Nothing is due

1:05

Tie out or throw out

Invoice extraction, checked against a total the model did not produce

1:50

Numbers in, words out

Board pack commentary where your spreadsheet owns every figure

2:35

If it is not written down, it did not happen

The review step and the file note your auditor will accept

One number in this paragraph is not in the ledger

WHAT IT WROTE

Revenue rose €184k, 12% on August. Material costs rose €39k, broadly in line with volume. Personnel costs rose €27k, driven by the 3.1% collective wage increase in September. Freight out rose €12k.

WHAT IT WAS GIVEN

- Revenue: 1,712 against 1,528
- Material costs: 694 against 655
- Personnel costs: 598 against 571
- Freight out: 41 against 29

CHECKPOINT

Find it. Hands up the moment you have.

The 3.1% came from nowhere, and it reads perfectly

WHAT IT WROTE

Personnel costs rose €27k, driven by the 3.1% collective wage increase in September.

WHERE 3.1% CAME FROM

Nowhere. There was no wage figure in anything it was given. It is simply the most plausible sentence for a German manufacturer in September.

It fits perfectly, and that is the problem.

An invented figure reads exactly like a real one. Nothing about it looks wrong until somebody goes looking for its source.

Half of finance teams still need six days or more to close



The median has sat at about six days for a decade, through every tool that promised to fix it.

Ledge, Month end close benchmarks 2025, n = 100, vendor research. Median trend: APQC benchmarking, via Numeric, 2026.

01

Your newest preparer has never seen your ledger

By the end of this block you know what the model is doing when it writes your commentary, and where it is safe to start.

Every law firm has a junior who drafts everything and signs nothing

1

Drafts from the desk

Fast, well read, tireless. Knows only what is in the folder you handed over, not what is in the archive.

2

Cites a case that does not exist

Perfectly formatted, confidently argued, invented. It reads exactly like the real ones beside it.

3

Signs nothing

Nothing leaves under the firm's name until a partner has read it and put their initials on it.

Hold that picture for thirty seconds. The next slide is what it means for your close.

It writes what a plausible answer looks like, from what you pasted in



Hand over

Two trial balance columns and the drivers you know. It sees this and nothing else: not your ERP, not last year, not the contract.



Draft

It predicts the most plausible commentary for those figures. Plausible, not looked up, which is exactly how a 3.1% appears.



Sign

A person ties every number back to the source and puts a name on it. Until then it is a draft.

Hand over. Draft. Sign.

Every use of AI in this session has that heartbeat. The third step is the one that makes it finance work.

Most finance teams use it. Few trust it with the close.

56 %

Use AI in finance

Doubled since 2023, and still the lowest of any business function.

17 %

Use it in core workflows

The rest sit in pilots or on one person's laptop. 45% call it a limited pilot.

35 %

Use ChatGPT

The most used tool in finance is a general chat window, not a finance product.

CFO Connect with Oracle NetSuite, Remote and Spendesk, State of AI in Finance 2026.

Start where volume is high and a wrong draft is cheap to catch

MONTHLY VOLUME	Variance commentary, invoice fields Every month, every entity. A reviewer catches a bad draft in minutes	Board pack narrative High stakes and a lot of words. Draft it, then sign it line by line
	One off memos and lookups Cheap to get wrong, and too rare to be worth building for	Posting journals, releasing payments Rare and expensive. A person decides, and the model stays out of it

COST OF A WRONG DRAFT

Three things that never go into a chat window

1

Personal data

Salaries by name, customer IBANs, anything your DPO would call personal data.

2

Unreleased results

Figures the board has not seen yet, and anything under a disclosure rule.

3

Check the licence, not the logo

Your company contract for the tool is what covers you. A personal login on the same product is not covered.

That is why the rounds today run on our sample ledger.

Real figures go only into the tool your company licensed, and only with the names taken out.

Three ways in, and every one of them is real work

A • RUN IT

Your company tool

Your approved AI tool and a real trial balance extract, names removed. Only if your licence covers it.

B • LITE

Our sample ledger

Any chat window and the sample we hand over. The same exercise, and the same lesson.

C • WATCH

Follow on screen

Watch this round, then run it at your next close with the sheet we hand over.

CHECKPOINT

Call out your letter. A, B or C.

Write the commentary, then make it show its working

TYPE THIS

Write one sentence of variance commentary for the five largest movements. Then list every number you used and its row.

Paste August and September first, in thousands of euros. One rule: give it no reasons yet, and notice where it invents one.

CHECKPOINT

How many reasons in your commentary appear nowhere in what you pasted?

02

Tie out or throw out

By the end of this block every figure the model extracts ties to a total it did not produce.

The partner never rereads the brief. She checks the citations.

WHAT SHE SKIPS

Forty pages of argument. If every citation holds, the argument usually does too.


WHAT SHE CHECKS

Every case number, against the court record. One citation that does not exist and the whole brief goes back.

Check against something the junior did not write.

It is faster than rereading everything the junior did, and it finds what rereading misses.

Every extracted figure ties to a total the model did not produce



The model never checks its own work. A total it did not produce does.

Extraction pays for itself inside a year, if the total is yours

9 to 10 months

Payback on extraction, AP and AR

Across 160 finance leaders, surveyed January to April 2026.

3 to 6 months

Payback on reporting and variance

The fastest in a survey of 102 investor backed CFOs. Block 3 is about this one.

3

Never let it write the total row

Ask for rows only. A total the model adds up will agree with its own misreads, and so will a tired reviewer.

Gartner, September 2026, n = 160. Consero Global 2026 CFO Report, n = 102, investor backed firms only.


Two wrong entities can add up to one clean group



It surfaced in entity statutory reporting, not in the group close. Tie out per entity.

An illustrative example from Uniqus, Internal controls over generative AI, 2026: stale FX rates, 2,300 intercompany entries auto posted over three days. USD millions.

From 2028 your domestic invoices arrive as data, so aim extraction elsewhere



Wachstumschancengesetz, section 14 UStG; EDICOM, 2026; VATupdate, March 2026. Check the BMF circulars before you change a process.

The same three ways in, with invoices this time

A • RUN IT

Your own invoices

Five real supplier invoices from last month, in your company tool.

B • LITE

Our five invoices

The sample pack, including the one with the rounding trap in it.

C • WATCH

Watch, then Monday

Watch the tie out, then run it on Monday's supplier batch.

One rule this round: you add the total, not the model.

It comes from the invoices, in your spreadsheet, so it is free to disagree with the model.

CHECKPOINT

Hands up if your letter changed since Round 1.

Extract five invoices, then let the total catch it

TYPE THIS

```
Return one row per invoice: supplier,  
number, date, net, VAT rate, VAT,  
gross. Do not round. No total row.
```

Add the gross total yourself, from the invoices. If it differs from the sum of the rows, stop reading rows and find the one document it misread.

CHECKPOINT

Did your two totals agree? If not, which invoice broke them?

03

Numbers in, words out

By the end of this block you have a board pack paragraph whose every figure came from your spreadsheet.

A good brief quotes the contract. It never retypes it.

WHAT THE JUNIOR WRITES

The argument: why this clause matters, what it means for the client, what to do next.

WHAT THE JUNIOR NEVER WRITES

The figures and the dates. Those are copied from the exhibit, and the exhibit is attached.

The words belong to the junior. The numbers belong to the exhibit.

The model gets the table and the reasons. Your spreadsheet does the arithmetic.

The paragraph

The model writes it, from the two layers below and from nothing else

The reasons

What actually happened, from you: two new contracts, a price rise, a one off

The variance table

Computed in the spreadsheet: actual, prior, difference, percentage

The ledger

The source every figure in the paragraph traces back to

Every number in the paragraph is in the table. Every reason is in your notes.

Three checks that catch an invented figure without rereading a line

01

Tie it to a control total

Add up what the paragraph claims and compare it to the table total. Invented components almost never reconcile.

02

Take one input away

Delete one reason from your notes and run it again. The right answer says the reason is missing. The dangerous one writes a new reason.

03

Put it next to last period

An invented figure is plausible on its own and odd in a time series. Last month's paragraph is the fastest reviewer you have.

Adapted from Agentic, AI hallucinations in financial reports, 2026.

The same movement, and only one version you could sign

NUMBERS ONLY

What it guessed

"Revenue increased due to higher volumes and timing effects."

TABLE AND NOTES

What you can sign

"Revenue rose €184k: two new framework contracts, €131k, and the 1 September price rise, €53k."

3

Give it permission to leave a gap

Without a NOT EXPLAINED line in the prompt, it fills every gap with a reason that sounds right.

Every figure in the middle card is in the table. Every reason is in your notes.

Never let it do the **arithmetic**.

It writes the sentence. The spreadsheet owns the number. Asked to calculate and explain in one breath, it does both confidently, and only one of them is its job.

Last round, and the one closest to your real month

A • RUN IT

Your own variance table

Last month against the month before, computed in your spreadsheet, names removed.

B • LITE

Our sample table

The Stuttgart sample, with three reasons already written out for you.

C • WATCH

Watch, then your close

Watch the three checks run, then use them on your next board pack.

CHECKPOINT

Who is bringing a real table? Hands up.

Write the board pack paragraph from your own table

TYPE THIS

Write the board pack paragraph from this table and my notes. Use only its numbers. If unexplained, write NOT EXPLAINED.

Paste the finished table and three things that really happened. One rule: no arithmetic by the model. Then run the three checks.

CHECKPOINT

Did it write NOT EXPLAINED anywhere? Read that line out.

04

If it is not written down, it did not happen

Every law firm keeps a file note: who decided what, on which document, and when. By the end of this block you can write yours for the AI step.

GoBD has no AI chapter. It does not need one.

PRINCIPLE	IN THE GOBD	FOR AN AI DRAFTED ENTRY
Traceability	Nachvollziehbarkeit	Who checked which draft, against what
Verifiability	Nachprüfbarkeit	Source and draft are both kept
Immutability	Unveränderbarkeit	Approved version logged, not overwritten
Completeness	Vollständigkeit	The AI step is in your process documentation

The principles already cover every preparer. The model is simply a new one.

GoBD, BMF letter; Lexware on Verfahrensdokumentation, 2026. Not legal advice: agree the wording with your Steuerberater.

Two people and a log stand between the draft and the ledger

	DRAFT	TIE OUT	REVIEW	POST
Preparer	Runs the prompt in the approved tool	Adds the total, marks the exceptions		
Reviewer		Opens two exceptions at random	Reads every figure against the table, initials it	Releases it
The log	Keeps the input and the draft		Records who approved, and when	Keeps the posted version

Your close already has a four eyes principle. The AI step goes through it, not around it.

One paragraph in your process documentation covers the AI step

1

What the tool does

Drafts variance commentary and extracts invoice fields. It posts nothing.

2

What a person checks

Ties every figure to the variance table or a control total, then initials the draft.

3

Write it for the process, once

Not per entry. Describe the step, the check and the log, and update it when the tool changes.

That paragraph is your file note. It is the first thing an auditor reads.

What I would not trust it with yet

01

Judging intent

Accrual, timing or control failure is a judgement about your business. It will offer one, fluently.

02

The productivity numbers are softer than the surveys

Executives reported AI gains averaging 1.8% in 2025, and revenue data implies much less, possibly a one year lag. Only 17% of CFOs see significant value from productivity AI.

03

Much of the extraction work is going away in Germany

From 2028 domestic invoices arrive structured. Build for contracts and foreign suppliers instead.

Duke Fuqua with the Federal Reserve Banks of Richmond and Atlanta, March 2026, about 750 executives. Gartner, July 2026, n = 204.

Your auditor will never ask which model you used. They will ask **who checked it.**

The tool will change twice before your next audit. The tie out, the table and the file note will still be there, with a name against each one.

What goes back to your desk on Monday

1

Keep the three checks

Tie to a total, take an input away, compare to last period. They work on any tool and any preparer.

2

Keep the rule

Numbers in, words out. The spreadsheet calculates, the model writes, and NOT EXPLAINED beats a guess.

3

Do this at your next close

Run one AI drafted variance paragraph through the three checks. Send me the first number that did not tie.

Nobody on your team will announce they are using it.

The commentary is already being drafted in a browser tab, and nobody has written down who checked it.

